



## II-Improve DNN - L3

- Hyperparameter tuning
  - Tuning process
    - The importance of hp (hyperparameters) - Ng's opinion
      - most IM
        - learning\_rate
      - second most IM
        - momentum  $\beta$
        - #hidden units
        - mini-batch size
      - third most IM
        - #hidden layers
        - learning\_rate decay
      - use default
        - Adam parameters:  $\beta_1$ ,  $\beta_2$ ,  $\epsilon$  (0.9, 0.999,  $10^{-8}$ )
    - how to try out the hp
      - random instead of grid
        - random gives more try for each of the hp



- coarse to fine
  - shrink the range of potentials

Coarse to fine



- Using an appropriate scale to pick hyperparameters
  - A
- Hyperparameters tuning in practice: Pandas vs. Caviar
  - A
- Batch Normalization
  - Normalizing activations in a network
    - A
  - Fitting Batch Norm into a neural network
    - A
  - Why does Batch Norm work?
    - A
  - Batch Norm at test time
    - A
- Multi-class classification
  - Softmax Regression
    - A
  - Training a softmax classifier
    - A
- Introduction to programming frameworks
  - Deep learning frameworks
    - Why framework
      - to code in scratch is to understand the detail about the algorithm, to know how the benethe goes
      - but for large applications and complex NNs, better to use framework in order to make use of their efficient computation and algorithm.

- popular frameworks
  - Caffe/Caffe2, CNTK, DL4J, Keras, Lasagne, mxnet, PaddlePaddle, TensorFlow, Theano, Torch
  - to do: research for their pros & cons
- choosing criteria
  - ease of programming both for development & deployment
  - running speed
  - truly open - open source with good governance
    - not maintained by single company or group
  - objective
    - application: computer vision, NLP, online advertising....
    - language
- TensorFlow
  - basic structure:
    - **Writing and running programs in TensorFlow has the following steps:**
      - **Create Tensors (variables) that are not yet executed/evaluated.**
      - **Write operations between those Tensors.**
        - **to construct the computation graph, no computation here.**
      - **Initialize your Tensors.**

- **Create a Session.**

**Method 1:**

```
sess = tf.Session()
# Run the variables initialization (if needed), run the operations
result = sess.run(..., feed_dict = {...})
sess.close() # Close the session
```

**Method 2:**

```
with tf.Session() as sess:
    # run the variables initialization (if needed), run the operations
    result = sess.run(..., feed_dict = {...})
    # This takes care of closing the session for you :)
```

- **Run the Session. This will run the operations you'd written above.**

- **only this step executes the computation**

- e.g.

- `w = tf.Variable(value, dtype)`
  - specify the training variables
- `x = tf.placeholder(tf.float32, [3,1])`
  - specify a variable whose value can be assigned later which is convenient to feed data later e.g. mini\_batches
  - to feed the data, use `feed_dict={x:coefficients, ...}`
- `cost = #function()`
  - define the cost function
- `train =`  
`tf.train.XXXOptimizer(learning_rate).minimize(cost)`
  - define the optimization algorithm
- `init = tf.global_variables_initializer()`
- `sess = tf.Session()`
- `sess.run(init)`
  - all has this
  - to start a session, another better way
    - with `tf.Session()` as `sess`
      - `sess.run(init)`

•  
`print(sess.run(train  
)`

- `sess.run(train, feed_dict={x:coefficients})` #run the optimization once

- to run more, put in into a for loop

- `print(sess.run(w))`

- basic comments

- the cost function is the key which build the **computation graph**, e.g. forward prop
- there's no need to worry about back prob, since TF takes care of it automatically.

- Ex - TF

- package

- `h5py`
- `matplotlib`
- `from .. import`

- TensorFlow - basic

- `y = tf.constant(36, name='y')`

- `tf.constant(value, dtype=None, shape=None, name='Const', verify_shape=False)`

- related

- `tf.zeros(shape, dtype=tf.float32, name=None)`
- `tf.range(start, limit=None, delta=1, dtype=None, name='range')`
- `tf.random_normal(shape, mean=0.0, stddev=1.0, dtype=tf.float32, seed=None, name=None)`

- `loss = tf.Variable((y - y_hat)**2, name='loss')`

- meaning

- A variable maintains state in the graph across calls to `run()`. You add a variable to the graph by constructing an instance of the class `Variable`.

- implement

- The `Variable()` constructor requires an initial value for the variable

- The initial value defines and fixes the type and shape of the variable

- The value can be changed using one of the assign methods

- to change the shape of a variable later you have to use an `assign Op` with

- `validate_shape=False`, e.g.

- `w.assign(w + 1.0)`

- `global_variables()` returns the contents of the graph collection

- `GraphKeys.GLOBAL_VARIABLES`

- which automatically collects variables into the graph

- program in TF

- create tensors **#that are not yet executed/evaluated.**

- build operation among tensors

- initialize tensor

- create a session
- run the operations inside the session #  
evaluate the value of all variables
- placeholder
  - `tf.placeholder(dtype, shape=None, name=None)`
  - A placeholder is an object whose value you can specify only later. To specify values for a placeholder, you can pass in values by using a "feed dictionary" (`feed_dict = {var:x,...}` variable).
- Variable vs. placeholder
  - property
    - variable: fixed shape & type; initialization;
    - placeholder: flexible shape & type; no ;
  - function
    - variable: to store state in the graph; initialization;
      - build the computation graph
    - placeholder: to input external data; no ;
      - pass parameter to function
      - pass training data
- `tf.one_hot(labels, depth, axis)`
  - Conversion: single number to vector. e.g. 3 => [0,0,1,0...]
- `tf.ones(shape, dtype = tf.float32, name = None)`
  - `tf.zeros()`

- `tf.ones_like(tensor, dtype = None, name = None, optimize=True)`

- Given a single tensor (`tensor`), this operation returns a tensor of the same type and shape as `tensor` with all elements set to 1.

## ■ COST

- `tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits = Z3, labels = Y))`

- **the "logits" and "labels" of `tf.nn.softmax_cross_entropy_with_logits` are expected to be of shape (number of examples, num\_classes).**

## ◦ numpy

- `b = a.reshape(a.shape[0], -1)` #keep the first dim, the rest are flattened to be a vector.

- `a = np.random.randn(3, 4, 2, 2)`
- `b.shape = (3, 16)`

- `train = train.astype(np.float32)`

- **`ndarray.astype`**(*dtype, order='K', casting='unsafe', subok=True, copy=True*)
- Copy of the array, cast to a specified type. Return a array with specific type