

Python - Quickstart tutorial for Numpy

- 2017/11/17 23:05
- <https://docs.scipy.org/doc/numpy-dev/user/quickstart.html>
- numpy
 - the basics
 - NumPy's main object is the **homogeneous multidimensional array**. It is a table of elements (usually numbers), all of the same type, indexed by a **tuple of positive integers**. In NumPy dimensions are called *axes*. The number of axes is *rank*.
 - NumPy's array class is called `ndarray`.
 - the Standard Python Library class `array.array`, which only handles one-dimensional arrays and offers less functionality
 - important attributes of an `ndarray` object are
 - `ndarray.ndim`
 - the number of axes (dimensions) of the array. `rank`
 - `ndarray.shape`
 - the dimensions of the array. `(n,m)`.
 - `ndarray.size`
 - the total number of elements of the array.
 - `ndarray.dtype`
 - an object describing the type of the elements in the array.
 - One can create or specify dtype's using standard Python types. Additionally NumPy provides types of its own. `numpy.int32`, `numpy.int16`, and `numpy.float64` are some examples.
 - `ndarray.itemsize`
 - the size in bytes of each element of the array.
 - It is equivalent to `ndarray.dtype.itemsize`.
 - `ndarray.data`
 - the buffer containing the actual elements of the array.
 - Creation
 - create an array from a regular Python list or tuple using the `array` function
 - 1D: `a = np.array([2, 3, 4])`
 - 2D: `b = np.array([(1.5, 2, 3), (4, 5, 6)])`
 - `c = np.array([[1, 2], [3, 4]], dtype=complex)`
 - The type of the array can also be explicitly specified at creation time:
 - unknown content but known size
 - NumPy offers several functions to create arrays with initial placeholder content. These minimize

the necessity of growing arrays, an expensive operation.

- `np.zeros((3,4))`
- `np.ones((2,3,4), dtype=np.int16)`
- `np.empty((2,3))`
- By default, the dtype of the created array is `float64`.

■ To create sequences of numbers

- analogous to `range` in python
- **`numpy.arange`**(`[start,] stop, [step,] dtype=None`)

- When `arange` is used with floating point arguments, it is generally not possible to predict the number of elements obtained, due to the finite floating point precision

- so it is usually better to use the function `linspace`

- `np.linspace(0, 2, 9)`
- **`numpy.linspace`**(`start, stop, num=50, endpoint=True, retstep=False, dtype=None`)

- Return evenly spaced numbers over a specified interval.

- **`numpy.logspace`**(`start, stop, num=50, endpoint=True, base=10.0, dtype=None`)

- Return numbers spaced evenly on a log scale.

◦ Printing Arrays

- `print`
- 1D array is printed as row array

- Basic Operations

- Arithmetic operators on arrays apply *elementwise*. A new array is created and filled with the result.

- `+`, `-`, `/`, `%` ...

- `multiply`

- the product operator `*` operates elementwise in NumPy arrays

- The matrix product can be performed using the `dot` function or method:

- Some operations, such as `+=` and `*=`, act in place to modify an existing array rather than create a new one.

- When operating with arrays of different types, the type of the resulting array corresponds to the more general or precise one (a behavior known as upcasting).

- unary operations

- `a.sum()`, `min()`, `max()` ...

- By default, these operations apply to the array as though it were a list of numbers, regardless of its shape.

- However, by specifying the `axis` parameter you can apply an operation along the specified axis of an array:

- `axis = 0` along the 1st/row

- `axis = 1` along the 2nd/column

- Universal Functions

- such as `sin`, `cos`, and `exp`

- Within NumPy, these functions operate elementwise on an array, producing an array as output.

- Indexing, Slicing and Iterating

- One-dimensional arrays can be indexed, sliced and iterated over, much like [lists](#) and other Python sequences.

- `a[2:5]`

- `a[:6:2] = -1000`

- *# equivalent to `a[0:6:2] = -1000`; from start to position 6, exclusive, set every 2nd element to -1000*

- `a[: :-1]`

- *# reversed a*

- `for i in a:`

- iteration
 - Multidimensional arrays can have one index per axis. These indices are given in a tuple separated by commas:
 - almost the same as matlab
 - `b[2, 3]`
 - `b[0:5, 1]`
 - `b[:, 1]`
 - When fewer indices are provided than the number of axes, the missing indices are considered complete slices:
 - `b[-1]` *# the last row. Equivalent to `b[-1, :]`*
 - The expression within brackets in `b[i]` is treated as an `i` followed by as many instances of `:` as needed to represent the remaining axes. NumPy also allows you to write this using dots as `b[i, ...]`.
 - `x[1, 2, ...]` is equivalent to `x[1, 2, :, :, :]`,
 - Iterating over multidimensional arrays is done with respect to the first axis:

```
>>> for row in b:
...     print(row)
...
[0 1 2 3]
[10 11 12 13]
[20 21 22 23]
[30 31 32 33]
[40 41 42 43]
```

- if one wants to perform an operation on each element in the array, use `flat`
 - `for element in b.flat:`
- Shape Manipulation
 - Note that the following three commands all return a modified array, but do not change the original array:
 - `a.ravel()` *# returns the array, flattened*
 - “C-style”, that is, the rightmost index “changes the fastest”, so the element after `a[0, 0]` is `a[0, 1]`.
 - `a.reshape(6, 2)` *# returns the array with a modified shape*
 - `a.T` *# returns the array, transposed*
 - The `reshape` function returns its argument with a modified shape, whereas the `ndarray.resize` method modifies the array itself:

- If a dimension is given as -1 in a reshaping operation, the other dimensions are automatically calculated:

- `a.reshape(3, -1)`

- Stacking together different arrays

- For arrays of with more than two dimensions, [hstack](#) stacks along their second axes, [vstack](#) stacks along their first axes, and [concatenate](#) allows for an optional arguments giving the number of the axis along which the concatenation should happen.

- `np.vstack((a, b))`

- `np.hstack((a, b))`

- `np.r_[1:4, 0, 4]`

- `array([1, 2, 3, 0, 4])`

- Copies and Views

- No Copy at All

- Simple assignments make no copy of array objects or of their data.

- `>>> a = np.arange(12)`

- `>>> b = a` *# no new object is created*

- `>>> b is a` *# a and b are two names for the same ndarray object*

- *a & b are just names, referring to the same object*

- Python passes mutable objects as references, so function calls make no copy.

- call by reference

- View or Shallow Copy

- Different array objects can share the same data. The `view` method creates a new array object that looks at the same data.

- `>>> c = a.view()`

- `>>> c is a`

- `False`

- `>>> c.base is a` *# c is a view of the data owned by a*

- `True`

- `>>> c.shape = 2, 6` *# a's shape doesn't change*

```
>>> a.shape
(3, 4)
>>> c[0,4] = 1234 # a's data changes
```

- Slicing an array returns a view of it:

- `s[:] = 10` *# s[:] is a view of s. Note the difference between s=10 and s[:]=10*
- view can change the value but not the shape

- Deep Copy

- The `copy` method makes a complete copy of the array and its data.

- `>>> d = a.copy()` *# a new array object with new data is created*
- *# d doesn't share anything with a*

- Functions and Methods Overview

- Array Creation
- [arange](#), [array](#), [copy](#), [empty](#), [empty_like](#), [eye](#), [fromfile](#), [fromfunction](#), [identity](#), [linspace](#), [logspace](#), [mgrid](#), [ogrid](#), [ones](#), [ones_like](#), [r](#), [zeros](#), [zeros_like](#)
- Conversions
- [ndarray.astype](#), [atleast_1d](#), [atleast_2d](#), [atleast_3d](#), [mat](#)
- Manipulations
- [array_split](#), [column_stack](#), [concatenate](#), [diagonal](#), [dsplit](#), [dstack](#), [hsplit](#), [hstack](#), [ndarray.item](#), [newaxis](#), [ravel](#), [repeat](#), [reshape](#), [resize](#), [squeeze](#), [swapaxes](#), [take](#), [transpose](#), [vsplit](#), [vstack](#)
- Questions
- [all](#), [any](#), [nonzero](#), [where](#)
- Ordering
- [argmax](#), [argmin](#), [argsort](#), [max](#), [min](#), [ptp](#), [searchsorted](#), [sort](#)
- Operations
- [choose](#), [compress](#), [cumprod](#), [cumsum](#), [inner](#), [ndarray.fill](#), [imag](#), [prod](#), [put](#), [putmask](#), [real](#), [sum](#)
- Basic Statistics
- [cov](#), [mean](#), [std](#), [var](#)
- Basic Linear Algebra
- [cross](#), [dot](#), [outer](#), [linalg.svd](#), [vdot](#)

- Broadcasting rules

- 在dim = 1 的维度上，复制扩展到与最大的matrix 一样大小
- a “1” dim will be repeatedly prepended to the shapes of the smaller arrays until all the arrays have the same number of dimensions.
- The value of the array element is assumed to be the same along that dimension for the “broadcast” array.

- Fancy indexing and index tricks

- Indexing with Arrays of Indices

```

■ >>> a = np.arange(12)**2 # the first 12 square
numbers
>>> i = np.array( [ 1, 1, 3, 8, 5 ] ) # an array of indices
>>> a[i] # the elements of a at the
positions i
array([ 1, 1, 9, 64, 25])
>>>
>>> j = np.array( [ [ 3, 4], [ 9, 7 ] ] ) # a bidimensional array of
indices
>>> a[j] # the same shape as j
array([[ 9, 16],
       [81, 49]])

```

- 用array做index，维度按照index array来

■ When the indexed array `a` is multidimensional, a single array of indices refers to the first dimension of `a`. the other dim are :

■ Naturally, we can put `i` and `j` in a sequence (say a list) and then do the indexing with the list.

```

• >>> l = [i, j]
>>> a[l] #
equivalent to a[i, j]
array([[ 2, 5],
       [ 7, 11]])

```

■ However, we can not do this by putting `i` and `j` into an array, because this array will be interpreted as indexing the first dimension of `a`.

- list `[i, j]` and tuple `([i, j])` 可以看作多维index，但 `array([i, j])` 不能

■ You can also use indexing with arrays as a target to assign to:

- `a[[0, 0, 2]] = [1, 2, 3]`

◦ Indexing with Boolean Arrays

■ When we index arrays with arrays of (integer) indices we are providing the list of indices to pick. With boolean indices the approach is different; we explicitly choose which items in the array we want and which ones we don't.

■ use boolean arrays that have the same shape as the original array:

- `b = a > 4` # *b is a boolean with a's shape*
 - *b is a boolean array, with True where a[i] > 4*
- `a[b]` # *1d array with the selected elements*

- `c = a[b]`
- `a[b] = 0` *# All elements of 'a' higher than 4 become 0*

- The `ix_()` function

- The `ix_` function can be used to combine different vectors so as to obtain the result for each n-uplet.

- Linear Algebra

- See `linalg.py` in `numpy` folder for more

- Tricks and Tips

- “Automatic” Reshaping

- To change the dimensions of an array, you can omit one of the sizes which will then be deduced automatically:
- `a.shape = 2,-1,3` *# -1 means “whatever is needed”*

- Vector Stacking

- In NumPy this works via the functions `column_stack`, `dstack`, `hstack` and `vstack`, depending on the dimension in which the stacking is to be done.