

L5 - why object recognition is difficult

2017/11/08 22:27

- why object recognition is difficult
 - the machine process
 - to take a bunch of numbers that describe the intensities of pixels and go from there to the label of an object.
 - the difficulties
 - Segmentation: Real scenes are cluttered with other objects
 - which piece from which object?
 - hidden parts
 - Lighting
 - The intensities of the pixels are determined as much by the lighting as by the objects.
 - so same obj at different light often has different intensities
 - Deformation
 - Objects can deform in a variety of non-affine ways
 - e.g. various handwritten 2
 - Affordances
 - Object classes are often defined by how they are used
 - chairs is sth to sit on/in
 - viewpoint
 - Changes in viewpoint cause changes in images that standard learning methods cannot cope with.
 - the consequence
 - With so much uncertainties, it's gonna be very hard for any hand engineered program to be able to do a good job of those things.
- Ways to achieve viewpoint invariance
 - viewpoint variance
 - various viewpoints make the object show on different pixels in an image, which is unlike other ML problems

- Approach 1: The invariant feature approach

- How

- Extract a large, redundant set of features that are invariant under transformations

- Why

- With enough invariant features, there is only one way to assemble them into an object.
 - The relationships between features are captured by other features. With overlapping and redundant features, one feature will tell you how two other features are related.

- ATT

- But for recognition, we must avoid forming features from parts of different objects.

- then the difficulty goes to how to find this big feature set!

- Approach 2: The judicious normalization approach

- use a bounding box to set the coordinate of the object to avoid dimension-hopping problem.

- but if the box is required to be added by the machine, it's a chicken-egg problem. Manual work is fine but expensive

Put a box around the object and use it as a coordinate frame for a set of normalized pixels.


- This solves the dimension-hopping problem. If we choose the box correctly, the same part of an object always occurs on the same normalized pixels.

- The box can provide invariance to many degrees of freedom: translation, rotation, scale, shear, stretch...

But choosing the box is difficult because of:

- Segmentation errors, occlusion, unusual orientations.

We need to recognize the shape to get the box right!



We recognize this letter before we do mental rotation to decide if it's a mirror image.

- variant: The brute force normalization approach

- When training the recognizer, use well-segmented, upright images to fit the correct box.

- At test time try all possible boxes in a range of positions and scales.

- – This approach is widely used for detecting upright things like faces and house numbers in unsegmented images.

- – It is much more efficient if the recognizer can cope with some variation in position and scale

so that we can use a coarse grid when trying all possible boxes.

- **Approach 3:** Convolutional neural networks for hand-written digit recognition

- The replicated feature approach – CNN

- apply the same feature detectors with different position in an image to get one feature map

- Could also replicate across scale and orientation

- but it's tricky and expensive

- Replication greatly reduces the number of free parameters to be learned.

- the weights are shared

- Use several different feature types to get several feature maps

- Allows each patch of image to be represented in several ways.

- Backpropagation with weight constraints to train the nets

- We compute the gradients as usual, and then modify the gradients so that they satisfy the constraints, $w_i = w_j$.

To constrain: $w_1 = w_2$

we need: $\Delta w_1 = \Delta w_2$

compute: $\frac{\partial E}{\partial w_1}$ and $\frac{\partial E}{\partial w_2}$

use $\frac{\partial E}{\partial w_1} + \frac{\partial E}{\partial w_2}$ for w_1 and w_2

- What does replicating the feature detectors achieve

- it's equivariant translation not invariant translation

- the position of the obj changes, the neural activities change accordingly.

- **Invariant knowledge:** If a feature is useful in some locations during training, detectors for that feature will be available in all locations during testing.

- Pooling the outputs of replicated feature detectors

- Get a small amount of translational invariance at each level to give a single output to the next level.

- this reduce the #inputs to the next layer, thus allowing us to have many more different feature maps

- types

- averaging four neighboring replicated detectors
 - max, works slightly better

- Problem

- after several levels of pooling, we lost information about the precise position of the objects

- this makes it impossible to use the precise spatial relationships between high-level parts for recognition

- e.g. to detect the face works well, but to detect whose face may fail

- e.g. LeNet – handwritten digits recognizer

- Key character

- – Many hidden layers

- – Many maps of replicated units in each layer.

- – Pooling of the outputs of nearby replicated units.

- - A wide net that can cope with several characters at once even if they overlap.

- what's the theory behind it?

- - A clever way of training a complete system, not just a recognizer.

- max margin

- Look the impressive demos of LENET at [hHp://yann.lecun.com](http://yann.lecun.com)

- Priors and Prejudice - how to use the prior knowledge to train the NN

- one approach: use the knowledge to the design of NN, like LeNet

- factors

- local connectivity
- weight sharing
- activation functions

- This is less intrusive than handdesigning the features.

- - But it still prejudices the network towards the particular way of solving the problem that we had in mind.

- Alternatively, we can use our prior knowledge to

create a whole lot more training data

- it's intuitive. but it require a lot of work and long time to train

- and this can make a huge difference.

Ciresan et. al. create a many kinds of synthesis digit data, train only a dumb huge net and achieve only with 35 errors, even better than the LeNet at the error aspect.

- How to detect a significant drop in the error rate?
 - errors are a single number, telling not much about the goodness
 - The McNemar test tells

	model 1 wrong	model 1 right
model 2 wrong	29	1
model 2 right	11	9959

	model 1 wrong	model 1 right
model 2 wrong	15	15
model 2 right	25	9945

- on the left, it's a 11:1 ratio while on the right, it's a 5:3 ratio. So clearly, on the left case, model2 is much better while in the left case, model 1 is only slightly better

- Convolutional neural networks for object recognition
 - can the experience learned from MNIST be transferred to the 3D objects in the real world?
 - Recognizing real objects in color photographs is much more complicated than recognizing hand-written digits
 - - Hundred times as many classes (1000 vs 10)
 - - Hundred times as many pixels (256 x 256 color vs 28 x 28 gray)
 - - Two dimensional image of three-dimensional

scene.

- a lot of info has lost
- - Cluttered scenes requiring segmentation
 - overlapped, hidden...
 - the effect of light luminance
- - Multiple objects in each image.
- Will the same type of convolutional neural network work?
 - training a lot of synthesis images is still a problem for the current PC [2013]
 - currently, no clear answer
- the competition

The ILSVRC-2012 competition on ImageNet

- The dataset has 1.2 million high-resolution training images.
 - The classification task:
 - Get the “correct” class in your top 5 bets. There are 1000 classes.
 - The localization task:
 - For each bet, put a box around the object. Your box must have at least 50% overlap with the correct box.
 - Some of the best existing computer vision methods were tried on this dataset by leading computer vision groups from Oxford, INRIA, XRCE, ...
 - Computer vision systems use complicated multi-stage systems.
 - The early stages are typically hand-tuned by optimizing a few parameters.
- University of Toronto (Alex Krizhevsky) with only 16.4% error for the first bet
- some tricks used

- architecture

- 7 hidden layers not counting some max pooling layers.
- The early layers were convolutional.
- The last two layers were globally connected.

- The activation functions:

- ReLU in every hidden layer.

- These train much faster and are

more expressive
than logistic
units.

- **Competitive
normalization**

- to suppress
hidden activities
when nearby units
have stronger
activities. This
helps with
variations in
intensity.

- Data

- Train on random 224x224
patches from the 256x256
images to get more data.
Also use left-
right reflections of the
images.

- Regularization

- dropout

- This stops hidden units from relying too much on other hidden units.

- The hardware

- this is also a key factor

- He uses a very efficient implementation of convolutional nets on two Nvidia GTX 580 Graphics Processor Units (over 1000 fast little cores)
 - GPUs are very good for matrix-matrix multiplies.
 - GPUs have very high bandwidth to memory.
 - This allows him to train the network in a week.
 - It also makes it quick to combine results from 10 patches at test time.
- We can spread a network over many cores if we can communicate the states fast enough.
- As cores get cheaper and datasets get bigger, big neural nets will improve faster than old-fashioned (*i.e.* pre Oct 2012) computer vision systems.

- A2

- Key words in the description

- data set

- 4-grams

- The training set consists of 372,550 4-grams.

- 372550×4

- Each entry is an integer that is the index of a word in the vocabulary.

- The validation and test sets have 46,568 4-grams each.

- operation

- `> load data.mat`

- `> fieldnames(data)`

- `'data.vocab'`

- `'data.trainData', 'data.validData'`
and `'data.testData'`

- `> data.vocab`
 - To see the list of words in the vocabulary
- split the data
 - `>[train_x, train_t, valid_x, valid_t, test_x, test_t, vocab] = load_data(100);`
 - This will load the data, separate it into inputs and target, and make mini-batches of size 100 for the training set.
- train
 - `> model = train(1);`
 - This will train the model for one epoch
 - The training method will output a 'model' (a struct containing weights, biases and a list of words).

■ work

- You have to fill in parts of the code in `fprop.m` and `train.m`.
- Once the code is correctly filled-in, you will see that the cross entropy starts decreasing.
- At this point, try changing the hyperparameters (number of epochs, number of hidden units, learning rates, momentum, etc) and see what effect that has

on the training and validation cross entropy.